

padl

May 20, 2025

```
[18]: !pip install numpy==1.24.0 pandas scikit-learn torch torchvision matplotlib
      ↪scipy gensim
```

Collecting numpy==1.24.0

Downloading

numpy-1.24.0-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata  
(5.6 kB)

Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages  
(2.2.2)

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-  
packages (1.6.1)

Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages  
(2.6.0+cu124)

Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-  
packages (0.21.0+cu124)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-  
packages (3.10.0)

Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages  
(1.15.3)

Collecting gensim

Downloading

gensim-4.3.3-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata  
(8.1 kB)

Requirement already satisfied: python-dateutil>=2.8.2 in  
/usr/local/lib/python3.11/dist-packages (from pandas) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-  
packages (from pandas) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-  
packages (from pandas) (2025.2)

Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-  
packages (from scikit-learn) (1.5.0)

Requirement already satisfied: threadpoolctl>=3.1.0 in  
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.6.0)

Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-  
packages (from torch) (3.18.0)

Requirement already satisfied: typing-extensions>=4.10.0 in  
/usr/local/lib/python3.11/dist-packages (from torch) (4.13.2)

Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-

packages (from torch) (3.4.2)  
 Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages  
 (from torch) (3.1.6)  
 Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages  
 (from torch) (2025.3.2)  
 Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch)  
   Downloading nvidia\_cuda\_nvrtc\_cu12-12.4.127-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
 Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch)  
   Downloading nvidia\_cuda\_runtime\_cu12-12.4.127-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
 Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch)  
   Downloading nvidia\_cuda\_cupti\_cu12-12.4.127-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.6 kB)  
 Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch)  
   Downloading nvidia\_cudnn\_cu12-9.1.0.70-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.6 kB)  
 Collecting nvidia-cublas-cu12==12.4.5.8 (from torch)  
   Downloading nvidia\_cublas\_cu12-12.4.5.8-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
 Collecting nvidia-cufft-cu12==11.2.1.3 (from torch)  
   Downloading nvidia\_cufft\_cu12-11.2.1.3-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
 Collecting nvidia-curand-cu12==10.3.5.147 (from torch)  
   Downloading nvidia\_curand\_cu12-10.3.5.147-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
 Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch)  
   Downloading nvidia\_cusolver\_cu12-11.6.1.9-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.6 kB)  
 Collecting nvidia-cuspars-cu12==12.3.1.170 (from torch)  
   Downloading nvidia\_cuspars-cu12-12.3.1.170-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.6 kB)  
 Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in  
 /usr/local/lib/python3.11/dist-packages (from torch) (0.6.2)  
 Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in  
 /usr/local/lib/python3.11/dist-packages (from torch) (2.21.5)  
 Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in  
 /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)  
 Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch)  
   Downloading nvidia\_nvjitlink\_cu12-12.4.127-py3-none-  
 manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
 Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-  
 packages (from torch) (3.2.0)  
 Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-  
 packages (from torch) (1.13.1)  
 Requirement already satisfied: mpmath<1.4,>=1.1.0 in  
 /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch) (1.3.0)  
 Requirement already satisfied: pillow!=8.3.\*,>=5.3.0 in

```

/usr/local/lib/python3.11/dist-packages (from torchvision) (11.2.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.2)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-
packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (4.58.0)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.3)
Collecting scipy
  Downloading
scipy-1.13.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(60 kB)

60.6/60.6 kB
3.5 MB/s eta 0:00:00
Requirement already satisfied: smart-open>=1.8.1 in
/usr/local/lib/python3.11/dist-packages (from gensim) (7.1.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
Requirement already satisfied: wrapt in /usr/local/lib/python3.11/dist-packages
(from smart-open>=1.8.1->gensim) (1.17.2)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->torch) (3.0.2)
Downloading
numpy-1.24.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.3
MB)

17.3/17.3 MB
105.2 MB/s eta 0:00:00
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl
(363.4 MB)

363.4/363.4 MB
4.4 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (13.8 MB)

13.8/13.8 MB
114.4 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (24.6 MB)

24.6/24.6 MB
94.3 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (883 kB)

883.7/883.7 kB
52.5 MB/s eta 0:00:00

```

Downloading nvidia\_cudnn\_cu12-9.1.0.70-py3-none-manylinux2014\_x86\_64.whl  
(664.8 MB)

664.8/664.8 MB

2.1 MB/s eta 0:00:00

Downloading nvidia\_cufft\_cu12-11.2.1.3-py3-none-manylinux2014\_x86\_64.whl  
(211.5 MB)

211.5/211.5 MB

5.8 MB/s eta 0:00:00

Downloading nvidia\_curand\_cu12-10.3.5.147-py3-none-manylinux2014\_x86\_64.whl (56.3 MB)

56.3/56.3 MB

14.8 MB/s eta 0:00:00

Downloading nvidia\_cusolver\_cu12-11.6.1.9-py3-none-manylinux2014\_x86\_64.whl (127.9 MB)

127.9/127.9 MB

7.9 MB/s eta 0:00:00

Downloading nvidia\_cusparses\_cu12-12.3.1.170-py3-none-manylinux2014\_x86\_64.whl (207.5 MB)

207.5/207.5 MB

5.7 MB/s eta 0:00:00

Downloading nvidia\_nvjitlink\_cu12-12.4.127-py3-none-manylinux2014\_x86\_64.whl (21.1 MB)

21.1/21.1 MB

70.0 MB/s eta 0:00:00

Downloading  
gensim-4.3.3-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (26.7 MB)

26.7/26.7 MB

21.1 MB/s eta 0:00:00

Downloading  
scipy-1.13.1-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (38.6 MB)

38.6/38.6 MB

19.1 MB/s eta 0:00:00

Installing collected packages: nvidia-nvjitlink-cu12, nvidia-curand-cu12, nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-cuda-cupti-cu12, nvidia-cublas-cu12, numpy, scipy, nvidia-cusparses-cu12, nvidia-cudnn-cu12, nvidia-cusolver-cu12, gensim

Attempting uninstall: nvidia-nvjitlink-cu12

Found existing installation: nvidia-nvjitlink-cu12 12.5.82

Uninstalling nvidia-nvjitlink-cu12-12.5.82:

Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82

Attempting uninstall: nvidia-curand-cu12

Found existing installation: nvidia-curand-cu12 10.3.6.82

Uninstalling nvidia-curand-cu12-10.3.6.82:

Successfully uninstalled nvidia-curand-cu12-10.3.6.82

Attempting uninstall: nvidia-cufft-cu12

Found existing installation: nvidia-cufft-cu12 11.2.3.61

Uninstalling nvidia-cufft-cu12-11.2.3.61:  
    Successfully uninstalled nvidia-cufft-cu12-11.2.3.61  
Attempting uninstall: nvidia-cuda-runtime-cu12  
    Found existing installation: nvidia-cuda-runtime-cu12 12.5.82  
Uninstalling nvidia-cuda-runtime-cu12-12.5.82:  
    Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82  
Attempting uninstall: nvidia-cuda-nvrtc-cu12  
    Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82  
Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:  
    Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82  
Attempting uninstall: nvidia-cuda-cupti-cu12  
    Found existing installation: nvidia-cuda-cupti-cu12 12.5.82  
Uninstalling nvidia-cuda-cupti-cu12-12.5.82:  
    Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82  
Attempting uninstall: nvidia-cublas-cu12  
    Found existing installation: nvidia-cublas-cu12 12.5.3.2  
Uninstalling nvidia-cublas-cu12-12.5.3.2:  
    Successfully uninstalled nvidia-cublas-cu12-12.5.3.2  
Attempting uninstall: numpy  
    Found existing installation: numpy 2.0.2  
Uninstalling numpy-2.0.2:  
    Successfully uninstalled numpy-2.0.2  
Attempting uninstall: scipy  
    Found existing installation: scipy 1.15.3  
Uninstalling scipy-1.15.3:  
    Successfully uninstalled scipy-1.15.3  
Attempting uninstall: nvidia-cusparse-cu12  
    Found existing installation: nvidia-cusparse-cu12 12.5.1.3  
Uninstalling nvidia-cusparse-cu12-12.5.1.3:  
    Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3  
Attempting uninstall: nvidia-cudnn-cu12  
    Found existing installation: nvidia-cudnn-cu12 9.3.0.75  
Uninstalling nvidia-cudnn-cu12-9.3.0.75:  
    Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75  
Attempting uninstall: nvidia-cusolver-cu12  
    Found existing installation: nvidia-cusolver-cu12 11.6.3.83  
Uninstalling nvidia-cusolver-cu12-11.6.3.83:  
    Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

jax 0.5.2 requires numpy>=1.25, but you have numpy 1.24.0 which is incompatible.  
seaborn 0.13.2 requires numpy!=1.24.0,>=1.20, but you have numpy 1.24.0 which is incompatible.

tsfresh 0.21.0 requires scipy>=1.14.0; python\_version >= "3.10", but you have scipy 1.13.1 which is incompatible.

thinc 8.3.6 requires numpy<3.0.0,>=2.0.0, but you have numpy 1.24.0 which is incompatible.

treescop 0.1.9 requires numpy>=1.25.2, but you have numpy 1.24.0 which is incompatible.

pymc 5.22.0 requires numpy>=1.25.0, but you have numpy 1.24.0 which is incompatible.

chex 0.1.89 requires numpy>=1.24.1, but you have numpy 1.24.0 which is incompatible.

alumentations 2.0.6 requires numpy>=1.24.4, but you have numpy 1.24.0 which is incompatible.

blosc2 3.3.2 requires numpy>=1.26, but you have numpy 1.24.0 which is incompatible.

imbalanced-learn 0.13.0 requires numpy<3,>=1.24.3, but you have numpy 1.24.0 which is incompatible.

albucore 0.0.24 requires numpy>=1.24.4, but you have numpy 1.24.0 which is incompatible.

jaxlib 0.5.1 requires numpy>=1.25, but you have numpy 1.24.0 which is incompatible.

tensorflow 2.18.0 requires numpy<2.1.0,>=1.26.0, but you have numpy 1.24.0 which is incompatible.

Successfully installed gensim-4.3.3 numpy-1.24.0 nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cuspars-cu12-12.3.1.170 nvidia-nvjitlink-cu12-12.4.127 scipy-1.13.1

```
[80]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet,
↳RidgeCV
from sklearn.pipeline import make_pipeline, Pipeline
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score, confusion_matrix,
↳r2_score)
from gensim.models import Word2Vec
from scipy.stats import mode
```

##Question 1

Part (a)

```
[3]: # Load Q11 data
data = pd.read_csv("PADL-Q11-train.csv")
X = data.drop(columns=["out"])
y = data["out"]
```

```
[4]: # Split into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X, y, train_size=0.9,
↳random_state=123)
```

```
[5]: data.corr(method="pearson")["out"].drop("out")
```

```
[5]: X1    -0.022519
X2     0.487261
X3    -0.406628
X4     0.138193
X5     0.331039
Name: out, dtype: float64
```

```
[6]: alphas = [0.001, 0.01, 0.1, 1.0, 10.0, 100.0]
pipeline = make_pipeline(
    StandardScaler(),
    PolynomialFeatures(degree=2),
    RidgeCV(alphas=alphas, fit_intercept=True)
)
```

```
[7]: pipeline.fit(X_train, y_train)
r2 = r2_score(y_val, pipeline.predict(X_val))
```

```
[8]: print(f"Validation R²: {r2:.4f}")
      print(f"Best alpha: {pipeline.named_steps['ridgecv'].alpha_}")
```

Validation R²: 1.0000

Best alpha: 0.001

```
[9]: final_model = make_pipeline(
      StandardScaler(),
      PolynomialFeatures(degree=2),
      Ridge(alpha=0.001)
    )
```

```
[10]: final_model.fit(X, y)
```

```
[10]: Pipeline(steps=[('standardscaler', StandardScaler()),
                       ('polynomialfeatures', PolynomialFeatures()),
                       ('ridge', Ridge(alpha=0.001))])
```

```
[ ]: #Load test set
      test_data = pd.read_csv("PADL-Q11-unseen.csv")
      X_test = test_data.drop(columns=["out"])
      y_test = test_data["out"]

      y_pred = final_model.predict(X_test)
      r2 = r2_score(y_test, y_pred)
      print(f"R² on test set: {r2:.4f}")
```

Part (b)

```
[13]: # Load the training data
      data = pd.read_csv("PADL-Q12-train.csv")
      X = data.drop(columns=["out"])
      y = data["out"]
```

```
[14]: X.describe().T
```

```
[14]:
```

	count	mean	std	min	25%	50% \
X1	300.0	500.451643	294.131746	5.061584	242.004526	510.331875
X2	300.0	4.981790	2.964880	0.108377	2.461287	4.775509
X3	300.0	25.922546	14.634331	0.231601	13.082446	27.031299
X4	300.0	24.969496	14.731770	0.287933	12.045990	24.901573

	75%	max
X1	760.993690	990.505142
X2	7.633798	9.997177
X3	38.372979	49.896706
X4	36.863381	49.970686



```
[15]: #5-fold cross-validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)
```

```
[19]: # Baseline: Linear Regression
baseline = Pipeline([
    ("scaler", StandardScaler()),
    ("lr", LinearRegression())
])
lr_scores = cross_val_score(baseline, X, y, cv=kf, scoring="r2")
baseline_r2 = lr_scores.mean()
baseline.fit(X, y)
baseline_coefs = baseline.named_steps["lr"].coef_
baseline_coef_sum = np.sum(np.abs(baseline_coefs))
```

```
[20]: print("R2: {:.4f}".format(baseline_r2))
for i, c in enumerate(baseline_coefs):
    print("Feature {} coefficient: {:.4f}".format(i+1, c))
print("Sum of coefficients: {:.4f}".format(baseline_coef_sum))
```

R<sup>2</sup>: 0.9569

Feature 1 coefficient: 17.7428

Feature 2 coefficient: 8.9652

Feature 3 coefficient: 14.6941

Feature 4 coefficient: 1.1752

Sum of coefficients: 42.5774

The chosen regularisation method is the ElasticNet, which combines both L1 (Lasso) and L2 (Ridge) penalties.

```
[25]: # Regularisation: ElasticNet with CV over alpha
best_r2 = 0
best_alpha = None
best_coef_sum = float("inf")
alphas = np.logspace(-4, 2, 100)
best_model = None

for alpha in alphas:
    elastic_pipe = Pipeline([
        ("scaler", StandardScaler()),
        ("elastic_net", ElasticNet(alpha=alpha, l1_ratio=0.9, random_state=42))
    ])
    elastic_scores = cross_val_score(elastic_pipe, X, y, cv=kf, scoring="r2")
    r2 = elastic_scores.mean()
    elastic_pipe.fit(X, y)
    coefs = elastic_pipe.named_steps["elastic_net"].coef_
    coef_sum = np.sum(np.abs(coefs))

    if r2 >= 0.9 * baseline_r2 and coef_sum < best_coef_sum:
```

```

        best_r2 = r2
        best_alpha = alpha
        best_model = elastic_pipe
        best_coefs = coefs
        best_coef_sum = coef_sum

final_model = best_model

```

```

[26]: print("R²: {:.4f}".format(best_r2))
      print("Alpha: {:.4f}".format(best_alpha))
      for i, c in enumerate(best_coefs):
          print("Feature {} coefficient: {:.4f}".format(i+1, c))
      print("Sum of coefficients: {:.4f}".format(best_coef_sum))

```

```

R²: 0.8677
Alpha: 2.3101
Feature 1 coefficient: 12.7888
Feature 2 coefficient: 5.6027
Feature 3 coefficient: 10.3725
Feature 4 coefficient: 0.0000
Sum of coefficients: 28.7639

```

```

[27]: #Print the difference for each coefficient separately
      for i, (base_c, reg_c) in enumerate(zip(baseline_coefs, best_coefs)):
          diff = base_c - reg_c
          print(f"Feature {i+1}: {diff:.4f}")

```

```

Feature 1: 4.9541
Feature 2: 3.3626
Feature 3: 4.3216
Feature 4: 1.1752

```

The ElasticNet model reduced the  $R^2$  by 8.92%, while the coefficient magnitude was reduced by almost one-third(32.4%).

```

[ ]: #Load test set
      test_data = pd.read_csv("PADL-Q12-unseen.csv")
      X_test = test_data.drop(columns=["out"])
      y_test = test_data["out"]

      #Predict on test set
      y_pred = final_model.predict(X_test)
      r2 = r2_score(y_test, y_pred)
      print(f"R² on test set: {r2:.4f}")

```

Part (c)

```
[29]: #Load the data
data = pd.read_csv('PADL-Q13-train.csv')
X = data.drop(columns=["out"])
y = data["out"]
```

Dataset contains 300 samples and 5 features. This sample-to-feature ratio satisfies rule of thumb ( 10 samples per feature) and reduces the risk of overfitting

```
[30]: X.shape
```

```
[30]: (300, 5)
```

To obtain more reliable assessment of performance cross-validation was used. It provides better estimation how well model generalises to unseen data and detect potential overfitting. Since the dataset is quite small, cross-validation effectively uses all 300 samples for both training and validation, rather than creating single train/val split. Using 5-folds, each iteration trains on 240 samples and validates on 60.

```
[31]: # Set up 5-Fold Cross-Validation
kf = KFold(n_splits=5, shuffle=True, random_state=62)
```

```
[32]: baseline_pipeline = Pipeline([
    ("lr", LinearRegression())
])
```

```
[60]: scores = cross_val_score(baseline_pipeline, X, y, cv=kf, scoring='r2')
print(f"5-fold cross-validation R2: {scores.mean():.4f}")
```

5-fold cross-validation R<sup>2</sup>: 0.9613

```
[61]: baseline_pipeline.fit(X, y)
print("Baseline Coefficients:", baseline_pipeline.named_steps["lr"].coef_)
```

Baseline Coefficients: [ 0.28561068 0.4141668 0.18665374 -0.30949337  
0.0174621 ]

```
[62]: r2_score = baseline_pipeline.score(X, y)
print(f"Baseline R2: {r2_score:.4f}")
```

Baseline R<sup>2</sup>: 0.9659

The descriptive statistics show that all features are continuous and have no missing values, but have very different distributions.

```
[57]: X.describe().T
```

```
[57]:
```

	count	mean	std	min	25%	50%	75%	\
X1	300.0	-0.118030	1.653688	-5.208844	-1.204841	-0.110005	1.037322	
X2	300.0	0.183434	2.899741	-8.403038	-1.738648	0.065389	2.089494	

X3	300.0	0.392237	6.202502	-22.208877	-4.135754	0.139140	5.011779
X4	300.0	0.181127	3.080314	-12.520381	-1.933037	-0.060989	2.291150
X5	300.0	-0.296061	4.190317	-12.687886	-2.922096	-0.255402	2.503598

	max
X1	4.760530
X2	9.397246
X3	22.905032
X4	9.095478
X5	14.102427

Linear regression model already performs well on the raw data ( $R^2 = 96.31\%$ ). To further improve numerical stability and make sure that all features contribute equally, standard scaling was applied.

```
[63]: #Scaling the data
pipeline_scaled = Pipeline([
    ("scaler", StandardScaler()),
    ("lr", LinearRegression())
])
```

```
[64]: scores = cross_val_score(pipeline_scaled, X, y, cv=kf, scoring='r2')
print(f"5-fold cross-validation R2 with Scaling: {scores.mean():.4f} ")
```


5-fold cross-validation  $R^2$  with Scaling: 0.9613

```
[65]: pipeline_scaled.fit(X, y)
print("Scaled Coefficients:", pipeline_scaled.named_steps["lr"].coef_)
```

Scaled Coefficients: [ 0.47152317 1.19897318 1.15578906 -0.95174665 0.0730497 ]

```
[66]: #Inspecting features for multicollinearity
sns.heatmap(X.corr(), annot=True, cmap="coolwarm")
```

```
[66]: <Axes: >
```



Heatmap revealed that some of the features (X3 and X4, X2 and X4) have strong correlations. This is an indicator of multicollinearity, which can result in unstable and unreliable coefficient estimates. To address this issue, PCA is applied to transform features into uncorrelated components.

```
[67]: pipeline_scaled_pca = Pipeline([
      ("scaler", StandardScaler()),
      ("pca", PCA(n_components=0.99, random_state=42)),
      ("lr", LinearRegression())])
```

```
[68]: scores = cross_val_score(pipeline_scaled_pca, X, y, cv=kf, scoring='r2')
      print(f"5-fold cross-validation R2 with Scaling and PCA: {scores.mean():.4f}")
```

5-fold cross-validation R<sup>2</sup> with Scaling and PCA: 0.9616

PCA was applied with `n_components=0.99`, retaining 99% of variance to preserve the model performance.

```
[69]: pipeline_scaled_pca.fit(X, y)
      print("Scaled and PCA Coefficients:", pipeline_scaled_pca.named_steps["lr"].
      ↪coef_)
```

Scaled and PCA Coefficients: [-0.44648346 0.27848727 1.1149977 1.40977924]

```
[ ]: #Load test set
test_data = pd.read_csv("PADL-Q13-unseen.csv")
X_test = test_data.drop(columns=["out"])
y_test = test_data["out"]

y_pred = pipeline_scaled_pca.predict(X_test)
r2 = r2_score(y_test, y_pred)
print(f"R2 on test set: {r2:.4f}") # raw (non-cross-validated) R2 score
```

##Question 2

Part (a)

```
[71]: #Load data
data = pd.read_csv("PADL-Q2.csv")

X = data.drop(columns="y")
y = data["y"].values
```

```
[72]: n_clusters = data["y"].nunique()
print(f"Num of classes: {n_clusters}")
```

Num of classes: 4

```
[73]: #Scale the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
[74]: #Apply k-means++ clustering to original data
kmeans = KMeans(n_clusters=n_clusters, init='k-means++', n_init=50,
↳ random_state=43)
kmeans_labels = kmeans.fit_predict(X_scaled)
```

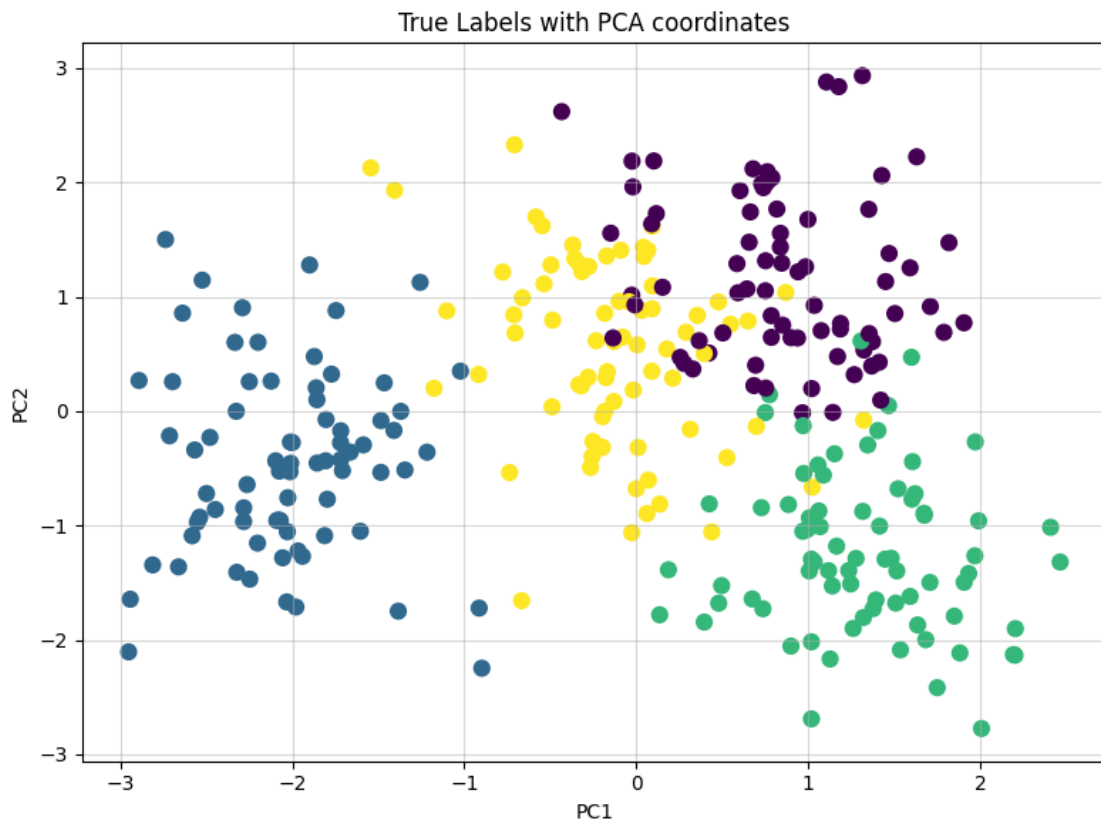
```
[75]: #Apply PCA with the first two components
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
```


```
[77]: def remap_labels(y_true, y_pred):
    result = np.zeros_like(y_pred)
    for i in np.unique(y_pred):
        mask = y_pred == i
        result[mask] = mode(y_true[mask], keepdims=False)[0]
    return result
```

```
[78]: mapped_labels = remap_labels(y, kmeans_labels)
acc = accuracy_score(y, mapped_labels)
```

```
[79]: plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', s=60)
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.title("True Labels with PCA coordinates")
plt.grid(True, alpha=0.5)
plt.tight_layout()
plt.show()

#K-means clusters with original data with PCA coordinates
plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=kmeans_labels, cmap='viridis', s=60)
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.title("K-means clustering with PCA coordinates")
plt.grid(True, alpha=0.5)
plt.tight_layout()
plt.show()
```






Part (b)

```
[82]: #KMeans on PCA-reduced data
kmeans_pca = KMeans(n_clusters=4, init='k-means++', n_init=50, random_state=42)
kmeans_labels_pca = kmeans_pca.fit_predict(X_pca)

[83]: mapped_labels_pca = remap_labels(y, kmeans_labels_pca)
acc_pca = accuracy_score(y, mapped_labels_pca)

[84]: plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=kmeans_labels_pca, cmap='viridis', s=60)
plt.xlabel("PC1");
plt.ylabel("PC2")
plt.title("K-means clustering on PCA-reduced data")
plt.grid(True, alpha=0.5)
plt.tight_layout()
plt.show()
```





Part (c)

```
[85]: explained_variance = pca.explained_variance_ratio_
print(f"Explained variance for PC1 and PC2 {explained_variance}")
print(f"Total variance explained {sum(explained_variance):.4f}")
```

Explained variance for PC1 and PC2 [0.37873093 0.2857641 ]  
Total variance explained 0.6645

```
[86]: print(f"Clustering accuracy: {acc * 100:.2f}%")
print(f"Clustering accuracy (PCA-reduced): {acc_pca * 100:.2f}%")
```

Clustering accuracy: 95.33%  
Clustering accuracy (PCA-reduced): 86.00%

```
[87]: #Print confusion matrices
cm = confusion_matrix(y, mapped_labels)
print("Confusion Matrix:")
print(cm)

cm_pca = confusion_matrix(y, mapped_labels_pca)
```

```
print("Confusion Matrix:")
print(cm_pca)
```

Confusion Matrix:

```
[[70  0  0  5]
 [ 0 75  0  0]
 [ 4  0 69  2]
 [ 2  1  0 72]]
```

Confusion Matrix:

```
[[59  0  0 16]
 [ 0 72  0  3]
 [ 6  0 69  0]
 [ 6  3  8 58]]
```

The relative loss of accuracy from using only two principal components is 9.33%, while the percentage of variance retained by PC1 and PC2 is 66.45%. This shows that although over one-third of the data variance was discarded, the clustering quality only dropped by ~10%, meaning that PC1 and PC2 capture most of the useful clustering structure.

##Question 3

Part (a): Calculate cosine similarities between node 5 and nodes 21-30

```
[ ]: with open("PADL-Q3.txt", "r") as f:
      walks = [line.strip().split() for line in f.readlines()]

model = Word2Vec(sentences=walks,
                  vector_size=64,
                  window=3,
                  workers=4,
                  min_count=1,
                  sg=1, #1 - Skip-gram
                  negative=15, #number of negative samples
                  epochs=30,
                  alpha=0.025, #initial learning rate
                  min_alpha=0.0001 #final learning rate
)

print("Cosine similarities between:")
for i in range(21, 31):
    sim = model.wv.similarity("5", str(i))
    print(f"node 5 and {i}: {sim:.4f}")
```

Cosine similarities between:

```
node 5 and 21 - 0.1768
node 5 and 22 - 0.1627
node 5 and 23 - 0.2832
node 5 and 24 - 0.3219
node 5 and 25 - 0.2156
```

node 5 and 26 - 0.2411  
node 5 and 27 - 0.2697  
node 5 and 28 - 0.2455  
node 5 and 29 - 0.2027  
node 5 and 30 - 0.2842

```
[ ]: num_walks = len(walks)
avg_walk_length = sum(len(walk) for walk in walks) / num_walks
max_walk_length = max(len(walk) for walk in walks)

print(f"Num of walks: {num_walks}")
print(f"Average walk length: {avg_walk_length}")
print(f"Max walk length: {max_walk_length}")
```

Num of walks - 5000  
Average walk length - 6.0  
Max walk length - 6

Part (b): Creating distance matrix sorted by similarity

```
[ ]: def create_distance_matrix(model):

    node_ids = sorted([int(node) for node in model.wv.index_to_key])
    distance_matrix = []

    for node in node_ids:
        node_str = str(node)

        similarities = [(str(other), model.wv.similarity(node_str, str(other)))
                        for other in node_ids if other != node]

        sorted_nodes = sorted(similarities, key=lambda x: x[1], reverse=True)
        row = node_str + ' ' + ' '.join(n for n, _ in sorted_nodes)
        distance_matrix.append(row)

    return distance_matrix, node_ids
```

```
[ ]: distance_matrix, node_ids = create_distance_matrix(model)
output_file = "PADL-Q3-result.txt"

with open(output_file, 'w') as f:
    for row in distance_matrix:
        f.write(row + "\n")
```

```
[ ]: model.wv.most_similar("5", topn=36)
```

```
[ ]: [('1', 0.8748427629470825),
      ('0', 0.874418318271637),
```

```
( '9', 0.8290271162986755),
( '6', 0.7461330890655518),
( '10', 0.7441762685775757),
( '11', 0.6938196420669556),
( '7', 0.6542657017707825),
( '14', 0.6419695019721985),
( '2', 0.6395747065544128),
( '3', 0.5698436498641968),
( '15', 0.5201167464256287),
( '19', 0.4299693703651428),
( '12', 0.4112282991409302),
( '8', 0.38661620020866394),
( '4', 0.38187888264656067),
( '18', 0.38183969259262085),
( '35', 0.3499222993850708),
( '16', 0.3462969660758972),
( '31', 0.3225603699684143),
( '24', 0.3218523859977722),
( '13', 0.3103030323982239),
( '30', 0.28419020771980286),
( '23', 0.28323984146118164),
( '20', 0.27852538228034973),
( '32', 0.27678021788597107),
( '27', 0.2696801424026489),
( '17', 0.26763835549354553),
( '28', 0.24553242325782776),
( '26', 0.2411223202943802),
( '33', 0.21819736063480377),
( '34', 0.21590673923492432),
( '25', 0.21555976569652557),
( '29', 0.20267580449581146),
( '21', 0.1767570823431015),
( '22', 0.16272933781147003)]
```

##Question 4

```
[ ]: import pandas as pd
import torch
import torch.nn as nn
import torch.optim as optim
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
```

```
[ ]: #Load and preprocess the data
data = pd.read_csv('body_measurements.csv')
data["Gender"] = data["Gender"].astype(int)
data = data.dropna()
```

```
[ ]: data.head()
```

```
[ ]:      Gender  Chest Circumference (mm)  Hip Circumference (mm)  Height (mm)  \
0         0          904.0          1000.0          1723.0
1         0          859.0          1027.0          1694.0
2         0          1092.0          1135.0          1659.0
3         1          1004.0          1091.0          2053.0
4         0          898.0          985.0          1608.0

      Weight (kg)  Waist Circumference (mm)
0         60.9          724.0
1         63.2          690.0
2         85.0          1014.0
3        107.6          916.0
4         61.3          755.0
```

Taking into account dataset characteristics (approximately 1800 samples with 5 input features), the model should be balancing complexity while still capturing relevant non-linear relationships.

```
[ ]: #Reorder to match the test set
X = data[['Chest Circumference (mm)', 'Hip Circumference (mm)', 'Height (mm)', 
↪ 'Weight (kg)', 'Gender']].values
y = data['Waist Circumference (mm)'].values.reshape(-1, 1)

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, 
↪ random_state=42)

scaler_X = StandardScaler()
scaler_y = StandardScaler()

X_train_scaled = scaler_X.fit_transform(X_train)
X_val_scaled = scaler_X.transform(X_val)

y_train_scaled = scaler_y.fit_transform(y_train)
y_val_scaled = scaler_y.transform(y_val)

X_train_tensor = torch.FloatTensor(X_train_scaled)
y_train_tensor = torch.FloatTensor(y_train_scaled)
X_val_tensor = torch.FloatTensor(X_val_scaled)
y_val_tensor = torch.FloatTensor(y_val_scaled)
```

```
[ ]: #Get feature-wise means and stds to normalise inputs in predict_waist.py
X_means = scaler_X.mean_
X_stds = scaler_X.scale_
y_mean = scaler_y.mean_
y_std = scaler_y.scale_

print("feature means:", [float(f"{m:.9f}") for m in X_means])
print("feature stds:", [float(f"{s:.9f}") for s in X_stds])
print("target mean: ", float(f"{y_mean[0]:.9f}"))
print("target std;", float(f"{y_std[0]:.9f}"))
```

```
feature means: [982.398174157, 1023.808988764, 1721.676966292, 74.307303371,
0.481039326]
feature stds: [109.335458295, 84.948100666, 106.074476316, 16.190130075,
0.499640363]
target mean: 854.990168539
target std; 120.186412625
```

```
[ ]: corr_with_waist = data.corr(numeric_only=True)["Waist Circumference (mm)"].
    ↪sort_values(ascending=False)
print("Correlations with Waist Circumference:")
print(corr_with_waist)
```

```
Correlations with Waist Circumference:
Waist Circumference (mm)    1.000000
Weight (kg)                 0.909266
Chest Circumference (mm)    0.892092
Hip Circumference (mm)      0.708293
Height (mm)                 0.366675
Gender                     0.301083
Name: Waist Circumference (mm), dtype: float64
```

The created network is a moderately deep feedforward neural network, designed for a regression task. It balances its capacity with generalisation ability given the dataset size of 1800 samples and small input dimensionality. The model has 2 hidden layers with 64 and 16 units respectively, followed by batch normalisation, ReLU activation and dropout rate of 0.2.

```
[ ]: class PredictWaist(nn.Module):
    def __init__(self, input_size=5):
        super(PredictWaist, self).__init__()

        self.layer1 = nn.Sequential(
            nn.Linear(input_size, 64),
            nn.BatchNorm1d(64),
            nn.ReLU(),
            nn.Dropout(0.2)
        )
```

```

self.layer2 = nn.Sequential(
    nn.Linear(64, 16),
    nn.BatchNorm1d(16),
    nn.ReLU(),
    nn.Dropout(0.2)
)

self.output_layer = nn.Linear(16, 1)

def forward(self, x):
    x = self.layer1(x)
    x = self.layer2(x)
    x = self.output_layer(x)
    return x

```

Huber Loss was used for its robustness to outliers and Adam optimizer was selected for faster and stable convergence. The model was kept compact to prevent overfitting, as more complex architectures did not lead to any performance gains.

```

[ ]: model = PredictWaist()
      criterion = nn.HuberLoss(delta=1.0)
      optimizer = optim.Adam(model.parameters(), lr=0.01)

```

```

[ ]: num_epochs = 100
      train_losses = []
      val_losses = []

      for epoch in range(num_epochs):
          model.train()
          optimizer.zero_grad()
          outputs = model(X_train_tensor)
          loss = criterion(outputs, y_train_tensor)
          loss.backward()
          optimizer.step()
          train_losses.append(loss.item())

          model.eval()
          with torch.no_grad():
              val_outputs = model(X_val_tensor)
              val_loss = criterion(val_outputs, y_val_tensor)
              val_losses.append(val_loss.item())

          if (epoch + 1) % 10 == 0:
              print(f'Epoch {epoch+1} Training loss - {loss.item():.4f}, Validation_
↳ loss - {val_loss.item():.4f}')

```

```

Epoch 10 Training loss - 0.0987, Validation loss - 0.1132
Epoch 20 Training loss - 0.0899, Validation loss - 0.0637

```

```
Epoch 30 Training loss - 0.0843, Validation loss - 0.0639
Epoch 40 Training loss - 0.0826, Validation loss - 0.0608
Epoch 50 Training loss - 0.0843, Validation loss - 0.0605
Epoch 60 Training loss - 0.0839, Validation loss - 0.0605
Epoch 70 Training loss - 0.0829, Validation loss - 0.0606
Epoch 80 Training loss - 0.0823, Validation loss - 0.0605
Epoch 90 Training loss - 0.0782, Validation loss - 0.0611
Epoch 100 Training loss - 0.0810, Validation loss - 0.0606
```

The validation loss decreased steadily, but plateau after epoch 50.

```
[ ]: model.eval()
with torch.no_grad():
    val_predictions_scaled = model(X_val_tensor)
    val_predictions = scaler_y.inverse_transform(val_predictions_scaled.numpy())

    abs_errors = np.abs(val_predictions - y_val)

    mae = np.mean(abs_errors)
    mape = np.mean(abs_errors / y_val) * 100
    print(f'Validation MAE: {mae:.2f} mm')
    print(f'Validation MAPE: {mape:.2f}%')

    top_10_errors = np.argsort(abs_errors.flatten())[::-1][:10]

    feature_names = ['Chest Circumference (mm)', 'Hip Circumference (mm)',
                     'Height (mm)', 'Weight (kg)', 'Gender']

    print("\nTop 10 Largest Errors:")
    for i in top_10_errors:
        actual = y_val[i][0]
        predicted = val_predictions[i][0]
        error = abs_errors[i][0]
        features = X_val[i]
        features_str = ', '.join(f"{name}={value}" for name, value in
    ↪ zip(feature_names, features))
        print(f"Index {i} | Error = {error:.2f} mm | Pred = {predicted:.2f} mm,
    ↪ True = {actual:.2f} mm")
        print(f"{features_str}")
```

Validation MAE: 32.49 mm

Validation MAPE: 3.86%

Top 10 Largest Errors:


```
Index 307 | Error = 162.60 mm | Pred = 1256.40 mm | True = 1419.00 mm
Chest Circumference (mm)=1301.0, Hip Circumference (mm)=1255.0, Height
(mm)=1934.0, Weight (kg)=146.6, Gender=1.0
Index 321 | Error = 128.64 mm | Pred = 1281.36 mm | True = 1410.00 mm
```



Chest Circumference (mm)=1282.0, Hip Circumference (mm)=1525.0, Height (mm)=1794.0, Weight (kg)=142.7, Gender=1.0  
 Index 99 | Error = 127.36 mm | Pred = 809.64 mm | True = 937.00 mm  
 Chest Circumference (mm)=984.0, Hip Circumference (mm)=1021.0, Height (mm)=1626.0, Weight (kg)=66.4, Gender=0.0  
 Index 220 | Error = 127.15 mm | Pred = 832.85 mm | True = 960.00 mm  
 Chest Circumference (mm)=938.0, Hip Circumference (mm)=938.0, Height (mm)=1679.0, Weight (kg)=70.0, Gender=1.0  
 Index 51 | Error = 124.08 mm | Pred = 698.92 mm | True = 823.00 mm  
 Chest Circumference (mm)=819.0, Hip Circumference (mm)=896.0, Height (mm)=1664.0, Weight (kg)=51.8, Gender=0.0  
 Index 139 | Error = 118.12 mm | Pred = 1027.88 mm | True = 1146.00 mm  
 Chest Circumference (mm)=1137.0, Hip Circumference (mm)=1063.0, Height (mm)=1617.0, Weight (kg)=90.1, Gender=1.0  
 Index 251 | Error = 112.55 mm | Pred = 760.45 mm | True = 873.00 mm  
 Chest Circumference (mm)=892.0, Hip Circumference (mm)=1049.0, Height (mm)=1630.0, Weight (kg)=62.1, Gender=0.0  
 Index 254 | Error = 110.82 mm | Pred = 901.18 mm | True = 1012.00 mm  
 Chest Circumference (mm)=1030.0, Hip Circumference (mm)=1087.0, Height (mm)=1586.0, Weight (kg)=77.4, Gender=0.0  
 Index 62 | Error = 100.04 mm | Pred = 883.96 mm | True = 984.00 mm  
 Chest Circumference (mm)=1015.0, Hip Circumference (mm)=1078.0, Height (mm)=1606.0, Weight (kg)=76.5, Gender=0.0  
 Index 67 | Error = 98.53 mm | Pred = 770.53 mm | True = 672.00 mm  
 Chest Circumference (mm)=943.0, Hip Circumference (mm)=971.0, Height (mm)=1708.0, Weight (kg)=62.5, Gender=0.0

The error analysis shows that challenges primarily occur at the extremes of the measurement distribution or for individuals with unusual body proportions. To deal with this, deeper architectures, feature engineering and upsampling of these underrepresented samples were tested, but none resulted in noticeable performance improvements. The model's performance plateaued, suggesting that the input features may have inherent limitations in predicting waist measurements beyond the achieved threshold. Nevertheless, the model achieved a mean absolute error (MAE) of approximately 32.5 mm and a mean absolute percentage error (MAPE) of 3.86%, indicating that predictions were still relatively accurate within about 4% of the actual waist measurements.

```
[ ]: plt.figure(figsize=(10, 6))
plt.scatter(y_val, val_predictions)
plt.xlabel('Actual Waist ')
plt.ylabel('Predicted Waist')
plt.title('Actual vs. Predicted Waist Values')
plt.plot([min(y_val), max(y_val)], [min(y_val), max(y_val)], 'r-')
plt.show()
```



```
[ ]: torch.save(model.state_dict(), "q4_model.pth")
```

##Question 5

```
[ ]: import random
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, Subset
from torchvision import transforms
import matplotlib.pyplot as plt
from PIL import Image
import numpy as np
import zipfile
import shutil
```

```
[ ]: with zipfile.ZipFile("garment_images.zip", 'r') as zip_ref:
    zip_ref.extractall(".")

os.mkdir("garment_images")

for cls in ['0', '1', '2']:
    if os.path.exists(cls):
```

```
shutil.move(cls, os.path.join("garment_images", cls))
```

```
[ ]: class GarmentDataset(Dataset):
    def __init__(self, root_dir, transform=None):
        self.root_dir = root_dir
        self.transform = transform
        self.images = []
        self.labels = []

        for class_name in ['0', '1', '2']:
            class_dir = os.path.join(root_dir, class_name)
            for img_name in os.listdir(class_dir):
                self.images.append(os.path.join(class_dir, img_name))
                self.labels.append(int(class_name))

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        img_path = self.images[idx]
        image = Image.open(img_path)
        label = self.labels[idx]
        if self.transform:
            image = self.transform(image)
        return image, label
```

```
[ ]: root_dir = "garment_images"
class_names = ['0', '1', '2']

fig, axs = plt.subplots(1, 3, figsize=(12, 4))
for i, class_name in enumerate(class_names):
    class_dir = os.path.join(root_dir, class_name)
    img_name = next(f for f in os.listdir(class_dir))
    img_path = os.path.join(class_dir, img_name)
    img = Image.open(img_path)
    axs[i].imshow(img)
    axs[i].set_title(f"Class {class_name}")
    axs[i].axis("off")
plt.tight_layout()
plt.show()
```



Only normalisation is applied. Data augmentation was also tested but it led to worse performance on validation set.

```
[ ]: data_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
↪225]))])

[ ]: full_dataset = GarmentDataset(root_dir="garment_images",
↪transform=data_transform)

[ ]: indices = list(range(len(full_dataset)))
    random.shuffle(indices)

    train_size = int(0.9 * len(full_dataset))
    train_indices = indices[:train_size]
    val_indices = indices[train_size:]
    val_size = len(val_indices)

    train_dataset = Subset(full_dataset, train_indices)
    val_dataset = Subset(full_dataset, val_indices)

[ ]: train_loader = DataLoader(dataset=train_dataset, batch_size=64, shuffle=True)
    val_loader = DataLoader(dataset=val_dataset, batch_size=64, shuffle=False)

[ ]: print(f"Total dataset: {len(full_dataset)}")
    print(f"Training set: {train_size}")
    print(f"Validation set: {val_size}")
    class_counts = [0, 0, 0]
    for _, label in full_dataset:
        class_counts[label] += 1
    print(f"T-shirts: {class_counts[0]}, Jumpers/Hoodies: {class_counts[1]}, Jeans:
↪{class_counts[2]}")
```

Total dataset: 2627  
Training set: 2364  
Validation set: 263  
T-shirts: 1025, Jumpers/Hoodies: 907, Jeans: 695

The dataset is a bit imbalanced with jeans being underrepresented class. To address this class weights were applied inversely proportional to class frequencies in the loss function.

The convolutional neural network was inspired by VGG design. It uses stacked convolutional layers followed by batch normalisation, ReLU activation and max pooling for spatial downsampling. The first 7\*7 convolution with stride 2 is inspired by ResNet to get a larger receptive field early in the network. This is followed by 4 convolutional blocks with increased channel depth from 64 to 512 to progressively extract more abstract and complex visual features. These blocks were scaled to balance the parameter count and generalisation, remaining within memory limits. The feature extracting blocks are followed by an adaptive average pooling layer, that ensures a fixed-size 512-dimensional vector. The fully connected part of the network reduces this vector to 256, applying ReLU activation and then to the 3 garment classes. A dropout rate of 0.5 is applied before both transitions to prevent overfitting.

```
[ ]: class GarmentCNN(nn.Module):
    def __init__(self, num_classes=3):
        super(GarmentCNN, self).__init__()

        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Conv2d(128, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(128, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(256, 512, kernel_size=3, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(inplace=True),
            nn.Conv2d(512, 512, kernel_size=3, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(inplace=True),
```

```

        nn.MaxPool2d(kernel_size=2, stride=2),
    )

    self.classifier = nn.Sequential(
        nn.AdaptiveAvgPool2d((1, 1)),
        nn.Flatten(),
        nn.Dropout(0.5),
        nn.Linear(512, 256),
        nn.ReLU(inplace=True),
        nn.Dropout(0.5),
        nn.Linear(256, num_classes)
    )

    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x

```

```

[ ]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = GarmentCNN(num_classes=3).to(device)
total_params = sum(p.numel() for p in model.parameters())
print(f"total parameters: {total_params:,}") # check the number of parameters
↳ to make sure it fits to the size limitations

```

total parameters: 4,791,939

```

[ ]: print(model)

```

```

GarmentCNN(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3))
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (9): ReLU(inplace=True)
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (11): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

        (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (18): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (20): ReLU(inplace=True)
        (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (22): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (23): ReLU(inplace=True)
        (24): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (classifier): Sequential(
      (0): AdaptiveAvgPool2d(output_size=(1, 1))
      (1): Flatten(start_dim=1, end_dim=-1)
      (2): Dropout(p=0.5, inplace=False)
      (3): Linear(in_features=512, out_features=256, bias=True)
      (4): ReLU(inplace=True)
      (5): Dropout(p=0.5, inplace=False)
      (6): Linear(in_features=256, out_features=3, bias=True)
    )
  )
)

```

Since the model is solving a multi-class classification task, loss function should measure the divergence between predicted class probabilities and true class labels. For this task, CrossEntropyLoss was used with class weights to address the class imbalance that is present in the dataset. They are computed as the inverse of class frequency to give more importance to underrepresented classes. Adam optimizer with the initial learning rate of 0.0005 (reduced from default 0.001 for better stability) and weight decay 1e-4. A ReduceLROnPlateau scheduler to reduce learning rate if validation loss plateaued. Additionally, gradient clipping was added after backpropagation to prevent instability from large gradients.

```

[ ]: weights = [1/1025, 1/907, 1/695]
weights = torch.tensor(weights)
criterion = nn.CrossEntropyLoss(weight=weights.to(device))
optimizer = optim.Adam(model.parameters(), lr=0.0005, weight_decay=1e-4)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min',
↪factor=0.1, patience=5, verbose=True)

```

```
[ ]: def train_epoch(model, train_loader, criterion, optimizer):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        optimizer.step()

        running_loss += loss.item() * inputs.size(0)
        _, predicted = outputs.max(1)
        total += labels.size(0)
        correct += predicted.eq(labels).sum().item()

    epoch_loss = running_loss / total
    epoch_acc = 100 * correct / total

    return epoch_loss, epoch_acc
```

```
[ ]: def validate(model, val_loader, criterion, n_aug=5):
    model.eval()
    running_loss = 0.0
    correct = 0
    total = 0
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            loss = criterion(outputs, labels)

            class_correct = [0] * 3
            class_total = [0] * 3
            for label, pred in zip(all_labels, all_preds):
                class_total[label] += 1
                if label == pred:
```



```

        class_correct[label] += 1
        val_class_acc = [100 * c / t if t > 0 else 0 for c, t in
↪zip(class_correct, class_total)]
        running_loss += loss.item() * inputs.size(0)
        _, predicted = outputs.max(1)
        total += labels.size(0)
        correct += predicted.eq(labels).sum().item()

        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

    epoch_loss = running_loss / total
    epoch_acc = 100 * correct / total

    return epoch_loss, epoch_acc, all_preds, all_labels, val_class_acc

```

```

[ ]: num_epochs = 30
    train_losses = []
    train_accs = []
    val_losses = []
    val_accs = []
    val_class_accs = [[], [], []]
    best_val_acc = 0.0
    best_epoch = 0

    for epoch in range(num_epochs):
        train_loss, train_acc = train_epoch(model, train_loader, criterion,
↪optimizer)
        train_losses.append(train_loss)
        train_accs.append(train_acc)

        val_loss, val_acc, all_preds, all_labels, val_class_acc = validate(model,
↪val_loader, criterion)
        val_losses.append(val_loss)
        val_accs.append(val_acc)

        for i in range(3):
            val_class_accs[i].append(val_class_acc[i])
        scheduler.step(val_loss)

        print(f"Epoch {epoch+1}/{num_epochs}")
        print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.2f}%")
        print(f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%")
        print(f"T-shirt Acc: {val_class_acc[0]:.2f}%, Jumper/Hoody Acc:
↪{val_class_acc[1]:.2f}%, Jeans Acc: {val_class_acc[2]:.2f}%")

        if val_acc > best_val_acc:

```

```
best_val_acc = val_acc
best_epoch = epoch
torch.save(model.state_dict(), "q5_model.pth")
```

Epoch 1/30

Train Loss: 0.5823, Train Acc: 73.52%

Val Loss: 3.5928, Val Acc: 37.26%

T-shirt Acc: 1.05%, Jumper/Hoody Acc: 100.00%, Jeans Acc: 0.00%

Epoch 2/30

Train Loss: 0.2934, Train Acc: 88.79%

Val Loss: 2.9991, Val Acc: 49.81%

T-shirt Acc: 100.00%, Jumper/Hoody Acc: 21.28%, Jeans Acc: 19.40%

Epoch 3/30

Train Loss: 0.2456, Train Acc: 91.24%

Val Loss: 0.2234, Val Acc: 89.35%

T-shirt Acc: 92.63%, Jumper/Hoody Acc: 79.79%, Jeans Acc: 97.01%

Epoch 4/30

Train Loss: 0.1876, Train Acc: 93.10%

Val Loss: 0.4103, Val Acc: 89.73%

T-shirt Acc: 91.58%, Jumper/Hoody Acc: 97.87%, Jeans Acc: 77.61%

Epoch 5/30

Train Loss: 0.1937, Train Acc: 93.06%

Val Loss: 0.4699, Val Acc: 82.13%

T-shirt Acc: 58.95%, Jumper/Hoody Acc: 100.00%, Jeans Acc: 89.55%

Epoch 6/30

Train Loss: 0.1746, Train Acc: 93.06%

Val Loss: 0.3399, Val Acc: 87.07%

T-shirt Acc: 97.89%, Jumper/Hoody Acc: 65.96%, Jeans Acc: 100.00%

Epoch 7/30

Train Loss: 0.1617, Train Acc: 93.95%

Val Loss: 0.1711, Val Acc: 92.02%

T-shirt Acc: 80.00%, Jumper/Hoody Acc: 98.94%, Jeans Acc: 100.00%

Epoch 8/30

Train Loss: 0.1348, Train Acc: 94.92%

Val Loss: 0.2958, Val Acc: 87.07%

T-shirt Acc: 66.32%, Jumper/Hoody Acc: 98.94%, Jeans Acc: 100.00%

Epoch 9/30

Train Loss: 0.1277, Train Acc: 95.30%

Val Loss: 0.3538, Val Acc: 89.35%

T-shirt Acc: 96.84%, Jumper/Hoody Acc: 78.72%, Jeans Acc: 92.54%

Epoch 10/30

Train Loss: 0.1314, Train Acc: 95.26%

Val Loss: 0.1338, Val Acc: 93.54%

T-shirt Acc: 96.84%, Jumper/Hoody Acc: 86.17%, Jeans Acc: 100.00%

Epoch 11/30

Train Loss: 0.1067, Train Acc: 95.69%

Val Loss: 0.1394, Val Acc: 94.30%

T-shirt Acc: 93.68%, Jumper/Hoody Acc: 91.49%, Jeans Acc: 100.00%  
 Epoch 12/30  
 Train Loss: 0.1023, Train Acc: 96.19%  
 Val Loss: 0.2976, Val Acc: 87.45%  
 T-shirt Acc: 70.53%, Jumper/Hoody Acc: 96.81%, Jeans Acc: 98.51%  
 Epoch 13/30  
 Train Loss: 0.0955, Train Acc: 96.19%  
 Val Loss: 0.1540, Val Acc: 93.92%  
 T-shirt Acc: 97.89%, Jumper/Hoody Acc: 86.17%, Jeans Acc: 100.00%  
 Epoch 14/30  
 Train Loss: 0.1003, Train Acc: 96.15%  
 Val Loss: 0.4371, Val Acc: 86.69%  
 T-shirt Acc: 100.00%, Jumper/Hoody Acc: 67.02%, Jeans Acc: 95.52%  
 Epoch 15/30  
 Train Loss: 0.0786, Train Acc: 96.87%  
 Val Loss: 0.1126, Val Acc: 95.06%  
 T-shirt Acc: 95.79%, Jumper/Hoody Acc: 91.49%, Jeans Acc: 100.00%  
 Epoch 16/30  
 Train Loss: 0.0761, Train Acc: 96.95%  
 Val Loss: 0.3136, Val Acc: 92.02%  
 T-shirt Acc: 98.95%, Jumper/Hoody Acc: 82.98%, Jeans Acc: 94.03%  
 Epoch 17/30  
 Train Loss: 0.0747, Train Acc: 97.08%  
 Val Loss: 0.1910, Val Acc: 92.02%  
 T-shirt Acc: 84.21%, Jumper/Hoody Acc: 96.81%, Jeans Acc: 97.01%  
 Epoch 18/30  
 Train Loss: 0.0688, Train Acc: 97.76%  
 Val Loss: 0.4919, Val Acc: 79.47%  
 T-shirt Acc: 47.37%, Jumper/Hoody Acc: 96.81%, Jeans Acc: 100.00%  
 Epoch 19/30  
 Train Loss: 0.0587, Train Acc: 97.93%  
 Val Loss: 0.1135, Val Acc: 95.82%  
 T-shirt Acc: 95.79%, Jumper/Hoody Acc: 93.62%, Jeans Acc: 100.00%  
 Epoch 20/30  
 Train Loss: 0.0720, Train Acc: 97.21%  
 Val Loss: 0.3379, Val Acc: 88.21%  
 T-shirt Acc: 70.53%, Jumper/Hoody Acc: 97.87%, Jeans Acc: 100.00%  
 Epoch 21/30  
 Train Loss: 0.0840, Train Acc: 96.62%  
 Val Loss: 0.3149, Val Acc: 90.11%  
 T-shirt Acc: 85.26%, Jumper/Hoody Acc: 97.87%, Jeans Acc: 86.57%  
 Epoch 22/30  
 Train Loss: 0.0458, Train Acc: 98.18%  
 Val Loss: 0.1083, Val Acc: 96.58%  
 T-shirt Acc: 98.95%, Jumper/Hoody Acc: 92.55%, Jeans Acc: 100.00%  
 Epoch 23/30  
 Train Loss: 0.0310, Train Acc: 99.07%  
 Val Loss: 0.1117, Val Acc: 96.96%

T-shirt Acc: 98.95%, Jumper/Hoody Acc: 93.62%, Jeans Acc: 100.00%  
 Epoch 24/30  
 Train Loss: 0.0222, Train Acc: 99.20%  
 Val Loss: 0.1061, Val Acc: 96.58%  
 T-shirt Acc: 97.89%, Jumper/Hoody Acc: 93.62%, Jeans Acc: 100.00%  
 Epoch 25/30  
 Train Loss: 0.0189, Train Acc: 99.28%  
 Val Loss: 0.1178, Val Acc: 96.58%  
 T-shirt Acc: 98.95%, Jumper/Hoody Acc: 92.55%, Jeans Acc: 100.00%  
 Epoch 26/30  
 Train Loss: 0.0179, Train Acc: 99.53%  
 Val Loss: 0.1168, Val Acc: 96.96%  
 T-shirt Acc: 98.95%, Jumper/Hoody Acc: 93.62%, Jeans Acc: 100.00%  
 Epoch 27/30  
 Train Loss: 0.0172, Train Acc: 99.45%  
 Val Loss: 0.1081, Val Acc: 96.58%  
 T-shirt Acc: 97.89%, Jumper/Hoody Acc: 93.62%, Jeans Acc: 100.00%  
 Epoch 28/30  
 Train Loss: 0.0134, Train Acc: 99.66%  
 Val Loss: 0.1169, Val Acc: 96.20%  
 T-shirt Acc: 96.84%, Jumper/Hoody Acc: 93.62%, Jeans Acc: 100.00%  
 Epoch 29/30  
 Train Loss: 0.0128, Train Acc: 99.58%  
 Val Loss: 0.1092, Val Acc: 95.82%  
 T-shirt Acc: 96.84%, Jumper/Hoody Acc: 92.55%, Jeans Acc: 100.00%  
 Epoch 30/30  
 Train Loss: 0.0119, Train Acc: 99.62%  
 Val Loss: 0.1120, Val Acc: 96.96%  
 T-shirt Acc: 98.95%, Jumper/Hoody Acc: 93.62%, Jeans Acc: 100.00%

```
[ ]: model.load_state_dict(torch.load("q5_model.pth"))
      print(f"Loaded best model from epoch {best_epoch+1} with validation accuracy:␣
           ↳{best_val_acc:.2f}%")
```

Loaded best model from epoch 23 with validation accuracy: 96.96%

Even though jeans were the most underrepresented class, they were still the easiest to classify with striking accuracy of 100%. Jumpers/Hoodies had the most variability probably due to their visual similarity with T-shirts, but the model still achieved accuracy of around 94% for them as well.


```
[ ]: plt.figure(figsize=(12, 5))

      plt.subplot(1, 2, 1)
      plt.plot(train_losses, label='Train Loss')
      plt.plot(val_losses, label='Val Loss')
      plt.xlabel('Epoch')
      plt.ylabel('Loss')
      plt.title('Training and Validation Loss')
```

```
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(train_accs, label='Train Acc')
plt.plot(val_accs, label='Val Acc')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```



While training metrics remained stable throughout the entire training process, validation loss and accuracy had some fluctuations. This is due to the small validation set size of only 263 samples, which makes the metrics more sensitive to individual batch variations.

```
[ ]: model.eval()

images, labels = next(iter(val_loader))
images, labels = images.to(device), labels.to(device)

with torch.no_grad():
    outputs = model(images)
    preds = outputs.argmax(dim=1)

class_names = ['T-shirt', 'Jumper/Hoody', 'Jeans']
plt.figure(figsize=(12, 12))
for i in range(9):
    plt.subplot(3, 3, i + 1)
    true_label = class_names[labels[i].item()]
```

```

pred_label = class_names[preds[i].item()]
plt.title(f"True: {true_label}\nPred: {pred_label}", fontsize=12)
plt.axis("off")
img = images[i].cpu().permute(1, 2, 0).clamp(0, 1).numpy()
plt.imshow(img)

plt.tight_layout()
plt.show()

```



From the visualisation only one sample was misclassified where the garment appears sleeveless and visually reminds T-shirt. All other samples were classified correctly.

#Question 6

```
[1]: import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import numpy as np
from PIL import Image
import zipfile
import random
from skimage.metrics import structural_similarity as ssim
```

```
[2]: with zipfile.ZipFile("face_images.zip", 'r') as zip_ref:
    zip_ref.extractall(".")

os.makedirs("face_images", exist_ok=True)
for filename in os.listdir("."):
    if filename.endswith(".jpg"):
        os.rename(filename, os.path.join("face_images", filename))
```

```
[3]: class FaceImages(Dataset):
    def __init__(self, root_dir, transform=None):
        self.root_dir = root_dir
        self.transform = transform
        self.images = []
        for img_name in os.listdir(root_dir):
            self.images.append(os.path.join(root_dir, img_name))

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        img_path = self.images[idx]
        image = Image.open(img_path)

        if self.transform:
            image = self.transform(image)


        return image
```

```
[4]: transform = transforms.Compose([transforms.ToTensor()])
dataset = FaceImages(root_dir="face_images", transform=transform)
```

```
[5]: train_size = int(0.9 * len(dataset))
      val_size = len(dataset) - train_size
      train_dataset, val_dataset = torch.utils.data.random_split(dataset,
      ↪ [train_size, val_size])

[6]: train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
      val_loader = DataLoader(val_dataset, batch_size=64, shuffle=False)

[7]: img = dataset[678]
      plt.imshow(img.squeeze(), cmap='gray')
      plt.show()
```



The development of the model was an iterative process that began with a smaller model, carefully balancing performance, depth and model size. The final design is a convolutional autoencoder also inspired by VGG style, using convolutional blocks for feature extraction and reconstruction.

The encoder has five convolution layers, increasing feature maps from 32 to 384 to capture complex representations. The layers were kept from 32 -> 48 -> 96 -> 192 -> 384 rather than 32 -> 64 -> 128 -> 256 -> 512, as the it would exceed memory limits. The strided convolutions reduced the input dimentionfrom 192×60 to 6×5. This results in a flattened vector of 11520 units to balance compactness and expressiveness. After that it was reduced to 32D latent representation vector. The decoder mirrors the encoder's structure with transposed convolution to reconstruct the original image. LeakyReLU activations were used in the encoder,while in the decoder standard



ReLU was used. A Sigmoid activation is applied to ensure that intensity values are within (0,1) range. Batch normalisation was applied after each convolutional layer to stabilise learning.

```
[8]: class Autoencoder(nn.Module):
    def __init__(self, latent_dim=32):
        super().__init__()

        #Encoder
        self.encoder = nn.Sequential(
            nn.Conv2d(1, 32, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(32),
            nn.LeakyReLU(),
            nn.Conv2d(32, 48, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(48),
            nn.LeakyReLU(),
            nn.Conv2d(48, 96, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(96),
            nn.LeakyReLU(),
            nn.Conv2d(96, 192, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(192),
            nn.LeakyReLU(),
            nn.Conv2d(192, 384, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(384),
            nn.LeakyReLU()
        )
        self.flatten = nn.Flatten()
        self.fc_enc = nn.Linear(384 * 6 * 5, latent_dim)

        #Decoder
        self.fc_dec = nn.Sequential(
            nn.Linear(latent_dim, 384 * 6 * 5),
            nn.ReLU(0.1)
        )
        self.decoder = nn.Sequential(
            nn.Unflatten(1, (384, 6, 5)),
            nn.ConvTranspose2d(384, 192, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(192),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(192, 96, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(96),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(96, 48, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(48),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(48, 24, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(24),
            nn.ReLU(inplace=True),
```

```

        nn.ConvTranspose2d(24, 12, kernel_size=4, stride=2, padding=1),
        nn.BatchNorm2d(12),
        nn.ReLU(inplace=True),
        nn.Conv2d(12, 1, kernel_size=3, padding=1),
        nn.Sigmoid()
    )

    def forward(self, x):
        x = self.encoder(x)
        x = self.flatten(x)
        z = self.fc_enc(x)
        x = self.fc_dec(z)
        x = self.decoder(x)
        return x

```

```

[9]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = Autoencoder(latent_dim=32).to(device)
print(model)

```

```

Autoencoder(
  (encoder): Sequential(
    (0): Conv2d(1, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.01)
    (3): Conv2d(32, 48, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
    (4): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (5): LeakyReLU(negative_slope=0.01)
    (6): Conv2d(48, 96, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
    (7): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (8): LeakyReLU(negative_slope=0.01)
    (9): Conv2d(96, 192, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
    (10): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (11): LeakyReLU(negative_slope=0.01)
    (12): Conv2d(192, 384, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
    (13): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (14): LeakyReLU(negative_slope=0.01)
  )
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (fc_enc): Linear(in_features=11520, out_features=32, bias=True)
  (fc_dec): Sequential(
    (0): Linear(in_features=32, out_features=11520, bias=True)
    (1): ReLU(inplace=True)
  )
)

```

```

    )
    (decoder): Sequential(
      (0): Unflatten(dim=1, unflattened_size=(384, 6, 5))
      (1): ConvTranspose2d(384, 192, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1))
      (2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (3): ReLU(inplace=True)
      (4): ConvTranspose2d(192, 96, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
      (5): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (6): ReLU(inplace=True)
      (7): ConvTranspose2d(96, 48, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
      (8): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (9): ReLU(inplace=True)
      (10): ConvTranspose2d(48, 24, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
      (11): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (12): ReLU(inplace=True)
      (13): ConvTranspose2d(24, 12, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
      (14): BatchNorm2d(12, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (15): ReLU(inplace=True)
      (16): Conv2d(12, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (17): Sigmoid()
    )
  )
)

```

```

[10]: total_param = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"Total num of parameters: {total_param:,}") # check the number of
↳ parameters to make sure it fits to the size limitations

```

Total num of parameters: 3,897,017

The model was trained using the MSE loss and Adam optimizer. The initial learning rate reduced from 1e-3 to 5e-4 for better stability. The initially chosen scheduler ReduceLROnPlateau was replaced with a CosineAnnealingLR for a smoother learning rate decay, contributing to the final SSIM gains in the later stages of training.

```

[11]: criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=5e-4, weight_decay=1e-5 )
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=150)

```

```

[12]: best_val_ssim = 0.0
      train_losses = []
      val_losses = []
      val_ssims = []

      for epoch in range(150):
          model.train()
          train_loss = 0.0

          for batch_idx, images in enumerate(train_loader):
              images = images.to(device)

              reconstructed = model(images)
              loss = criterion(reconstructed, images)

              optimizer.zero_grad()
              loss.backward()
              optimizer.step()

              train_loss += loss.item()

          avg_train_loss = train_loss / len(train_loader)
          train_losses.append(avg_train_loss)

          model.eval()
          val_loss = 0.0
          val_ssim = 0.0
          val_images = 0

          with torch.no_grad():
              for images in val_loader:
                  images = images.to(device)
                  reconstructed = model(images)
                  loss = criterion(reconstructed, images)
                  val_loss += loss.item()

                  batch_ssim = 0.0
                  for i in range(images.size(0)):
                      img1 = images[i].squeeze().cpu().numpy()
                      img2 = reconstructed[i].squeeze().cpu().numpy()
                      batch_ssim += ssim(img1, img2, data_range=1.0)

                  val_ssim += batch_ssim
                  val_images += images.size(0)

          avg_val_loss = val_loss / len(val_loader)
          avg_val_ssim = val_ssim / val_images

```

```

val_losses.append(avg_val_loss)
val_ssims.append(avg_val_ssim)

print(f"Epoch {epoch+1}/150, Train Loss: {avg_train_loss:.4f}, Val Loss:␣
↪{avg_val_loss:.4f}, Val SSIM: {avg_val_ssim:.4f}")

scheduler.step()

if avg_val_ssim > best_val_ssim:
    best_val_ssim = avg_val_ssim
    best_epoch = epoch + 1
    torch.save(model.state_dict(), 'q6_model.pth')

print(f"Model saved from epoch {best_epoch} with val SSIM: {best_val_ssim:.4f}")

```

```

Epoch 1/150, Train Loss: 0.0374, Val Loss: 0.0173, Val SSIM: 0.4345
Epoch 2/150, Train Loss: 0.0149, Val Loss: 0.0134, Val SSIM: 0.4891
Epoch 3/150, Train Loss: 0.0112, Val Loss: 0.0097, Val SSIM: 0.5976
Epoch 4/150, Train Loss: 0.0097, Val Loss: 0.0081, Val SSIM: 0.5686
Epoch 5/150, Train Loss: 0.0084, Val Loss: 0.0077, Val SSIM: 0.6174
Epoch 6/150, Train Loss: 0.0070, Val Loss: 0.0063, Val SSIM: 0.6349
Epoch 7/150, Train Loss: 0.0061, Val Loss: 0.0062, Val SSIM: 0.6733
Epoch 8/150, Train Loss: 0.0057, Val Loss: 0.0051, Val SSIM: 0.6843
Epoch 9/150, Train Loss: 0.0053, Val Loss: 0.0060, Val SSIM: 0.6444
Epoch 10/150, Train Loss: 0.0050, Val Loss: 0.0047, Val SSIM: 0.7093
Epoch 11/150, Train Loss: 0.0046, Val Loss: 0.0059, Val SSIM: 0.6999
Epoch 12/150, Train Loss: 0.0045, Val Loss: 0.0047, Val SSIM: 0.7028
Epoch 13/150, Train Loss: 0.0045, Val Loss: 0.0046, Val SSIM: 0.6826
Epoch 14/150, Train Loss: 0.0041, Val Loss: 0.0037, Val SSIM: 0.7313
Epoch 15/150, Train Loss: 0.0041, Val Loss: 0.0039, Val SSIM: 0.7284
Epoch 16/150, Train Loss: 0.0037, Val Loss: 0.0039, Val SSIM: 0.7333
Epoch 17/150, Train Loss: 0.0039, Val Loss: 0.0036, Val SSIM: 0.7364
Epoch 18/150, Train Loss: 0.0033, Val Loss: 0.0035, Val SSIM: 0.7170
Epoch 19/150, Train Loss: 0.0031, Val Loss: 0.0069, Val SSIM: 0.6185
Epoch 20/150, Train Loss: 0.0030, Val Loss: 0.0036, Val SSIM: 0.7371
Epoch 21/150, Train Loss: 0.0030, Val Loss: 0.0031, Val SSIM: 0.7532
Epoch 22/150, Train Loss: 0.0032, Val Loss: 0.0032, Val SSIM: 0.7622
Epoch 23/150, Train Loss: 0.0030, Val Loss: 0.0032, Val SSIM: 0.7392
Epoch 24/150, Train Loss: 0.0027, Val Loss: 0.0038, Val SSIM: 0.7311
Epoch 25/150, Train Loss: 0.0030, Val Loss: 0.0051, Val SSIM: 0.7027
Epoch 26/150, Train Loss: 0.0031, Val Loss: 0.0047, Val SSIM: 0.7091
Epoch 27/150, Train Loss: 0.0028, Val Loss: 0.0039, Val SSIM: 0.6927
Epoch 28/150, Train Loss: 0.0028, Val Loss: 0.0031, Val SSIM: 0.7719
Epoch 29/150, Train Loss: 0.0027, Val Loss: 0.0032, Val SSIM: 0.7662
Epoch 30/150, Train Loss: 0.0025, Val Loss: 0.0028, Val SSIM: 0.7644
Epoch 31/150, Train Loss: 0.0027, Val Loss: 0.0033, Val SSIM: 0.7650
Epoch 32/150, Train Loss: 0.0025, Val Loss: 0.0030, Val SSIM: 0.7843

```


Epoch 33/150, Train Loss: 0.0025, Val Loss: 0.0046, Val SSIM: 0.7037  
Epoch 34/150, Train Loss: 0.0028, Val Loss: 0.0030, Val SSIM: 0.7351  
Epoch 35/150, Train Loss: 0.0025, Val Loss: 0.0029, Val SSIM: 0.7680  
Epoch 36/150, Train Loss: 0.0025, Val Loss: 0.0031, Val SSIM: 0.7689  
Epoch 37/150, Train Loss: 0.0024, Val Loss: 0.0029, Val SSIM: 0.7856  
Epoch 38/150, Train Loss: 0.0021, Val Loss: 0.0028, Val SSIM: 0.7617  
Epoch 39/150, Train Loss: 0.0023, Val Loss: 0.0028, Val SSIM: 0.7741  
Epoch 40/150, Train Loss: 0.0023, Val Loss: 0.0048, Val SSIM: 0.6934  
Epoch 41/150, Train Loss: 0.0021, Val Loss: 0.0052, Val SSIM: 0.6673  
Epoch 42/150, Train Loss: 0.0023, Val Loss: 0.0031, Val SSIM: 0.7324  
Epoch 43/150, Train Loss: 0.0021, Val Loss: 0.0029, Val SSIM: 0.7399  
Epoch 44/150, Train Loss: 0.0022, Val Loss: 0.0025, Val SSIM: 0.7985  
Epoch 45/150, Train Loss: 0.0023, Val Loss: 0.0030, Val SSIM: 0.7842  
Epoch 46/150, Train Loss: 0.0021, Val Loss: 0.0026, Val SSIM: 0.7924  
Epoch 47/150, Train Loss: 0.0020, Val Loss: 0.0025, Val SSIM: 0.7948  
Epoch 48/150, Train Loss: 0.0020, Val Loss: 0.0028, Val SSIM: 0.7863  
Epoch 49/150, Train Loss: 0.0021, Val Loss: 0.0038, Val SSIM: 0.7494  
Epoch 50/150, Train Loss: 0.0023, Val Loss: 0.0028, Val SSIM: 0.7939  
Epoch 51/150, Train Loss: 0.0022, Val Loss: 0.0025, Val SSIM: 0.7995  
Epoch 52/150, Train Loss: 0.0019, Val Loss: 0.0043, Val SSIM: 0.7229  
Epoch 53/150, Train Loss: 0.0020, Val Loss: 0.0037, Val SSIM: 0.7854  
Epoch 54/150, Train Loss: 0.0021, Val Loss: 0.0026, Val SSIM: 0.7448  
Epoch 55/150, Train Loss: 0.0019, Val Loss: 0.0025, Val SSIM: 0.8066  
Epoch 56/150, Train Loss: 0.0017, Val Loss: 0.0022, Val SSIM: 0.7970  
Epoch 57/150, Train Loss: 0.0020, Val Loss: 0.0030, Val SSIM: 0.8003  
Epoch 58/150, Train Loss: 0.0021, Val Loss: 0.0025, Val SSIM: 0.8061  
Epoch 59/150, Train Loss: 0.0017, Val Loss: 0.0023, Val SSIM: 0.8163  
Epoch 60/150, Train Loss: 0.0018, Val Loss: 0.0024, Val SSIM: 0.7996  
Epoch 61/150, Train Loss: 0.0017, Val Loss: 0.0023, Val SSIM: 0.8062  
Epoch 62/150, Train Loss: 0.0018, Val Loss: 0.0027, Val SSIM: 0.8035  
Epoch 63/150, Train Loss: 0.0018, Val Loss: 0.0024, Val SSIM: 0.8043  
Epoch 64/150, Train Loss: 0.0019, Val Loss: 0.0023, Val SSIM: 0.8103  
Epoch 65/150, Train Loss: 0.0017, Val Loss: 0.0025, Val SSIM: 0.7636  
Epoch 66/150, Train Loss: 0.0017, Val Loss: 0.0037, Val SSIM: 0.7440  
Epoch 67/150, Train Loss: 0.0019, Val Loss: 0.0023, Val SSIM: 0.8080  
Epoch 68/150, Train Loss: 0.0015, Val Loss: 0.0022, Val SSIM: 0.8128  
Epoch 69/150, Train Loss: 0.0016, Val Loss: 0.0025, Val SSIM: 0.7923  
Epoch 70/150, Train Loss: 0.0015, Val Loss: 0.0021, Val SSIM: 0.7900  
Epoch 71/150, Train Loss: 0.0016, Val Loss: 0.0020, Val SSIM: 0.8024  
Epoch 72/150, Train Loss: 0.0015, Val Loss: 0.0023, Val SSIM: 0.8165  
Epoch 73/150, Train Loss: 0.0015, Val Loss: 0.0020, Val SSIM: 0.8130  
Epoch 74/150, Train Loss: 0.0014, Val Loss: 0.0020, Val SSIM: 0.8276  
Epoch 75/150, Train Loss: 0.0014, Val Loss: 0.0022, Val SSIM: 0.8161  
Epoch 76/150, Train Loss: 0.0014, Val Loss: 0.0020, Val SSIM: 0.8247  
Epoch 77/150, Train Loss: 0.0016, Val Loss: 0.0021, Val SSIM: 0.8246  
Epoch 78/150, Train Loss: 0.0014, Val Loss: 0.0021, Val SSIM: 0.8272  
Epoch 79/150, Train Loss: 0.0015, Val Loss: 0.0022, Val SSIM: 0.8110  
Epoch 80/150, Train Loss: 0.0014, Val Loss: 0.0024, Val SSIM: 0.7983

Epoch 81/150, Train Loss: 0.0014, Val Loss: 0.0019, Val SSIM: 0.8238  
Epoch 82/150, Train Loss: 0.0013, Val Loss: 0.0020, Val SSIM: 0.8158  
Epoch 83/150, Train Loss: 0.0014, Val Loss: 0.0021, Val SSIM: 0.8235  
Epoch 84/150, Train Loss: 0.0014, Val Loss: 0.0020, Val SSIM: 0.8267  
Epoch 85/150, Train Loss: 0.0012, Val Loss: 0.0023, Val SSIM: 0.8125  
Epoch 86/150, Train Loss: 0.0013, Val Loss: 0.0020, Val SSIM: 0.8273  
Epoch 87/150, Train Loss: 0.0013, Val Loss: 0.0020, Val SSIM: 0.8215  
Epoch 88/150, Train Loss: 0.0012, Val Loss: 0.0026, Val SSIM: 0.7997  
Epoch 89/150, Train Loss: 0.0011, Val Loss: 0.0018, Val SSIM: 0.8350  
Epoch 90/150, Train Loss: 0.0011, Val Loss: 0.0019, Val SSIM: 0.8091  
Epoch 91/150, Train Loss: 0.0012, Val Loss: 0.0018, Val SSIM: 0.8336  
Epoch 92/150, Train Loss: 0.0010, Val Loss: 0.0018, Val SSIM: 0.8363  
Epoch 93/150, Train Loss: 0.0010, Val Loss: 0.0019, Val SSIM: 0.8243  
Epoch 94/150, Train Loss: 0.0011, Val Loss: 0.0022, Val SSIM: 0.8227  
Epoch 95/150, Train Loss: 0.0011, Val Loss: 0.0023, Val SSIM: 0.7961  
Epoch 96/150, Train Loss: 0.0012, Val Loss: 0.0018, Val SSIM: 0.8283  
Epoch 97/150, Train Loss: 0.0010, Val Loss: 0.0018, Val SSIM: 0.8394  
Epoch 98/150, Train Loss: 0.0010, Val Loss: 0.0018, Val SSIM: 0.8333  
Epoch 99/150, Train Loss: 0.0010, Val Loss: 0.0018, Val SSIM: 0.8310  
Epoch 100/150, Train Loss: 0.0011, Val Loss: 0.0018, Val SSIM: 0.8363  
Epoch 101/150, Train Loss: 0.0010, Val Loss: 0.0017, Val SSIM: 0.8329  
Epoch 102/150, Train Loss: 0.0009, Val Loss: 0.0017, Val SSIM: 0.8430  
Epoch 103/150, Train Loss: 0.0009, Val Loss: 0.0017, Val SSIM: 0.8453  
Epoch 104/150, Train Loss: 0.0009, Val Loss: 0.0017, Val SSIM: 0.8462  
Epoch 105/150, Train Loss: 0.0009, Val Loss: 0.0017, Val SSIM: 0.8415  
Epoch 106/150, Train Loss: 0.0009, Val Loss: 0.0017, Val SSIM: 0.8436  
Epoch 107/150, Train Loss: 0.0009, Val Loss: 0.0017, Val SSIM: 0.8445  
Epoch 108/150, Train Loss: 0.0009, Val Loss: 0.0016, Val SSIM: 0.8422  
Epoch 109/150, Train Loss: 0.0008, Val Loss: 0.0017, Val SSIM: 0.8420  
Epoch 110/150, Train Loss: 0.0008, Val Loss: 0.0021, Val SSIM: 0.8088  
Epoch 111/150, Train Loss: 0.0008, Val Loss: 0.0017, Val SSIM: 0.8416  
Epoch 112/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8474  
Epoch 113/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8508  
Epoch 114/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8502  
Epoch 115/150, Train Loss: 0.0008, Val Loss: 0.0016, Val SSIM: 0.8485  
Epoch 116/150, Train Loss: 0.0008, Val Loss: 0.0016, Val SSIM: 0.8500  
Epoch 117/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8419  
Epoch 118/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8499  
Epoch 119/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8489  
Epoch 120/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8449  
Epoch 121/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8464  
Epoch 122/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8487  
Epoch 123/150, Train Loss: 0.0007, Val Loss: 0.0015, Val SSIM: 0.8489  
Epoch 124/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8483  
Epoch 125/150, Train Loss: 0.0007, Val Loss: 0.0016, Val SSIM: 0.8525  
Epoch 126/150, Train Loss: 0.0007, Val Loss: 0.0015, Val SSIM: 0.8526  
Epoch 127/150, Train Loss: 0.0007, Val Loss: 0.0015, Val SSIM: 0.8468  
Epoch 128/150, Train Loss: 0.0006, Val Loss: 0.0016, Val SSIM: 0.8471

Epoch 129/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8551  
Epoch 130/150, Train Loss: 0.0006, Val Loss: 0.0016, Val SSIM: 0.8529  
Epoch 131/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8545  
Epoch 132/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8515  
Epoch 133/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8456  
Epoch 134/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8551  
Epoch 135/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8464  
Epoch 136/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8544  
Epoch 137/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8492  
Epoch 138/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8542  
Epoch 139/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8469  
Epoch 140/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8525  
Epoch 141/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8540  
Epoch 142/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8533  
Epoch 143/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8512  
Epoch 144/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8558  
Epoch 145/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8481  
Epoch 146/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8527  
Epoch 147/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8541  
Epoch 148/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8525  
Epoch 149/150, Train Loss: 0.0005, Val Loss: 0.0015, Val SSIM: 0.8505  
Epoch 150/150, Train Loss: 0.0006, Val Loss: 0.0015, Val SSIM: 0.8516  
Model saved from epoch 144 with val SSIM: 0.8558


```
[14]: plt.figure(figsize=(8, 5))  
plt.plot(train_losses, label='Train Loss')  
plt.plot(val_losses, label='Validation Loss')  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.title('Training and Validation Loss')  
plt.legend()  
plt.grid(True)  
plt.tight_layout()  
plt.show()
```





Judging by the consistent decrease in both training and validation losses, no overfitting occurred. The reduced learning rate of  $5e-4$  allowed more stable training, reducing initial fluctuations. The Adam optimizer contributed to improvements early on, while CosineAnnealingLR facilitated further gains in the later stages of training, preventing premature convergence. L2 regularisation (weight decay of  $1e-5$ ) helped to keep both curves aligned, indicating the model generalises well. The main gains occurred around 50 epochs, while training for the full 150 epochs allowed for additional incremental gains.

```
[15]: plt.figure(figsize=(8, 5))
plt.plot(val_ssims, label='Validation SSIM')
plt.xlabel('Epoch')
plt.ylabel('SSIM')
plt.title('Validation SSIM')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



Unlike the losses, SSIM experienced larger fluctuations, but showed consistent upward trend. The general trend stabilised after epoch 70, followed by smoother and consistent improvements.

```
[16]: images = next(iter(val_loader))
      images = images[-4:].to(device)


      with torch.no_grad():
          reconstructions = model(images)

      fig, axes = plt.subplots(2, 4, figsize=(15, 9))
      for i in range(4):
          orig = images[i].squeeze().cpu().numpy()
          recon = reconstructions[i].squeeze().cpu().numpy()
          score = ssim(orig, recon, data_range=1.0)

          axes[0, i].imshow(orig, cmap='gray')
          axes[0, i].axis('off')

          axes[1, i].imshow(recon, cmap='gray')
          axes[1, i].axis('off')
          axes[1, i].set_title(f"SSIM: {score:.4f}")

      plt.tight_layout()
      plt.show()
```



After visualisation of some reconstructions, model accurately reconstructed the main facial features, although some blurriness was present. SSIM scores are slightly lower for samples with poorer light conditions, indicating that model may underperform on dark regions.